University of Wollongong in Dubai

MAAI901 Deep Learning for Visual Data — Applied Vision Transformer Report

When Labels Are Scarce: Fine-Tuned versus Frozen CLIP Vision Transformers for Steel Surface Defect Classification

Karthik Krishna (ID 1697286) Instructor: Dr. Milan Dordevic 26 October 2026

Abstract

Steel plants rarely have large labelled defect archives. This report asks how many labelled images per class a vision transformer needs, and whether full fine-tuning is worth its cost over a frozen encoder with a linear classifier. On the NEU surface defect dataset (six classes, 1,800 images), frozen CLIP ViT-B/16 features, partial and full fine-tuning, an ImageNet-21k ViT-B/16 and zero-shot CLIP are compared, using 5-fold cross-validation, 1 to 240 labels per class, calibration, rare-defect, unseen-defect and corruption tests. Full fine-tuning of CLIP reaches 95% accuracy with 3 labels per class (tuned frozen features: 10) and leads a tuned frozen baseline by 12.7 points at 1 label, but the gap vanishes beyond about 50 labels. Fine-tuned models are well calibrated yet still make confident mistakes, and zero-shot CLIP performs at chance. Because the dataset saturates easily, the claims are limited accordingly.

1. Introduction

Automatic visual inspection of hot-rolled steel supports the UAE's goal of increasing its competitiveness by adopting artificial intelligence in manufacturing [1]. Labelling defects is an expert task, and defects are rare. Vision transformers (ViTs) [2] pretrained on large datasets transfer well, but users must choose between cheap frozen features and more expensive fine-tuning. This leads to the question: *how does the number of labelled images per class change the best way to adapt a pretrained ViT to steel defect classification, and what do accuracy numbers hide?* We therefore evaluate not only accuracy but also calibration, rare classes, unseen defects and image corruptions.

2. Dataset and Pre-processing

Dataset and justification. The study uses the NEU surface defect database (NEU-CLS) [3, 4]: 1,800 grayscale 200×200 images of hot-rolled steel strips, 300 for each of six defects (crazing, inclusion, patches, pitted surface, rolled-in scale, scratches). It is a published benchmark of a real industrial problem, balanced, small enough for repeated experiments on a laptop, and has class names that map to plant vocabulary. The images are stored as three-channel files; they are treated as grayscale content. A public Kaggle copy [5] was used; the official page asks users to cite the two papers above and states no formal licence, so both are cited and the data are used for coursework only.

Known limitations. The images come from one laboratory setup and are very homogeneous: every test image has a near-duplicate in the training pool (cosine similarity of CLIP embeddings: minimum 0.921, median 0.980; Fig. 5). The task therefore saturates, so the analysis concentrates on low label counts and on tests beyond accuracy. Results do not show how models perform on another plant, camera or steel grade. GC10-DET [6] (3,570 images, ten classes, closer to plant conditions) was considered, but its only confirmed download is a Baidu drive with an undocumented format, so NEU was kept and GC10-DET is listed as future work.

Splits. For each of 5 folds, 60 images per class are held out (360 test images; every image is tested exactly once across folds) and $k \in \{1, 2, 3, 5, 10, 20, 50, 240\}$ training images per class are drawn from the remaining 240 per class; $k \leq 20$ uses two random draws. Hyper-parameters were fixed before any test result was seen.

Table 1: Compared methods.

Method	Trained parameters	Notes
Zero-shot CLIP	none	text prompts chosen on the training pool
Frozen CLIP + LR	3k (linear head)	$C = 10$, or C tuned on validation images
Partial fine-tune	≈15M (last 2 of 12 blocks, norm, projection)	CLIP
Full fine-tune	≈86M	CLIP; ImageNet ViT as reference

The one tuned method (frozen classifier, regularisation chosen from six values) selects on 60 further images per class taken from the unused pool, never on test images.

Pre-processing. Images are resized to 224×224 (bicubic) and normalised with each backbone's own mean and standard deviation. Fine-tuning uses random-resized crops (scale 0.7–1.0), horizontal and vertical flips; evaluation uses no augmentation. No other filtering was applied.

3. Model Selection and Implementation

Backbones. An input image is split into non-overlapping chunks, called "patches", each of size 16×16 pixels. These patches are embedded and passed to 12 transformer blocks [2]. This paper adopts the ViT-B/16 architecture, with about 86M parameters, and two types of pretrained weights used as backbones. The first type, CLIP, is pretrained on image–text pairs [7] and uses the OpenAI weights from OpenCLIP [8], with QuickGELU as the activation. The motivation for using language supervision is that it gives broadly useful features and allows zero-shot testing. The second type is pretrained on ImageNet-21k and fine-tuned on ImageNet-1k [9], with the weights available through the timm library [10], and is used as a reference with different pretraining data.

Adaptation regimes (Table 1). Models are trained using cross-entropy loss, the AdamW optimiser, a batch size of 32, and at least 100 training steps (or 10 epochs for larger training sets). A cosine learning-rate schedule with 10% warm-up is used, with an encoder learning rate of 10^{-5} and a head learning rate of 10^{-3}. Mixed-precision training is used whenever possible. For rare-defect runs, inverse-frequency class weights are used, which is a common technique for imbalanced classes [11].

Implementation. We use PyTorch with OpenCLIP, timm and scikit-learn. The code, scripts and raw results are in the GitHub repository at https://github.com/kk917/neu-vit-steel-defects. The entire experimental run on an Apple M4 Air (MPS GPU) took about 27 hours of wall-clock time, including breaks. The experiments are: (1) learning curves; (2) a rare defect with $1, 5, 20$ labels while the other classes have 50 labels; (3) an unseen defect, training on five classes and evaluating on the sixth; (4) robustness to blur, noise, brightness and contrast, including a model trained with blur and noise augmentation; (5) sensitivity to learning rate and regularisation; and (6) zero-shot prompts.

Metrics. We report accuracy and macro-F1; log-loss and expected calibration error (ECE, 10 bins) [12]; and confident errors, the share of test images misclassified with probability ≥ 0.9. AUROC is used to detect unseen defects, with one minus the maximum softmax probability as the score [13]. Differences between methods are paired across folds and reported with 95% t-intervals ($n = 5$). The folds share training images, so the intervals are only approximate.

4. Evaluation and Performance Analysis

Learning curves. Fig. 1 and Table 2 show full fine-tuning of CLIP is best at every budget. Against frozen features with tuned regularisation it leads by 12.7 points at $k = 1$ (95% CI 8.4 to 17.1), 8.1 at $k = 2$, 4.5 at

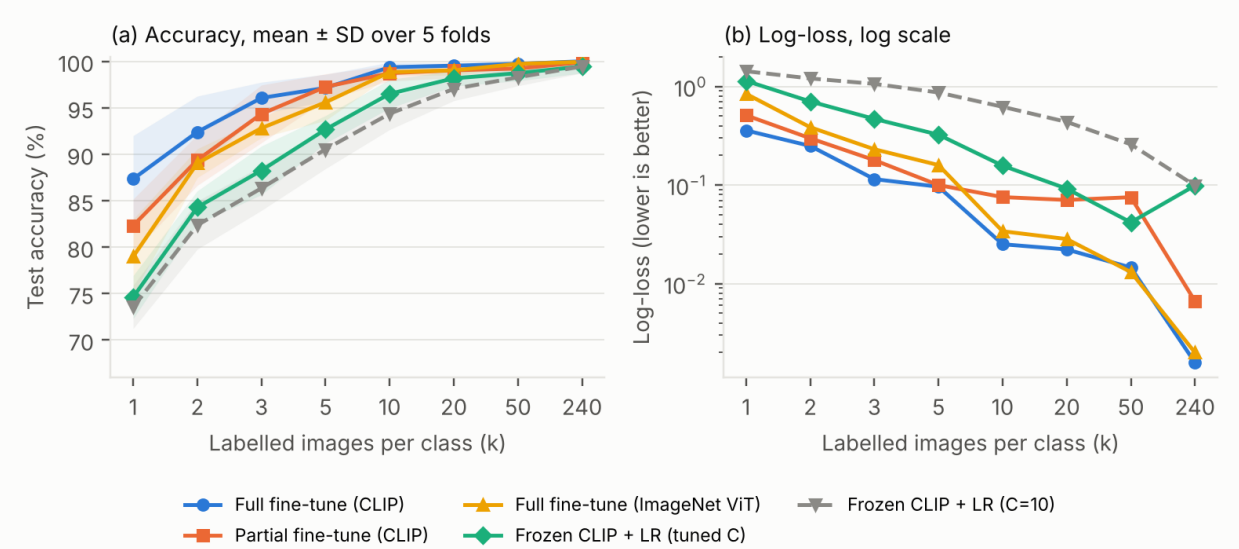


Figure 1: Learning curves (5 folds). (a) Accuracy, shaded ±1 SD over folds. (b) Log-loss.

Table 2: Test accuracy (%), mean ± SD over 5 folds.

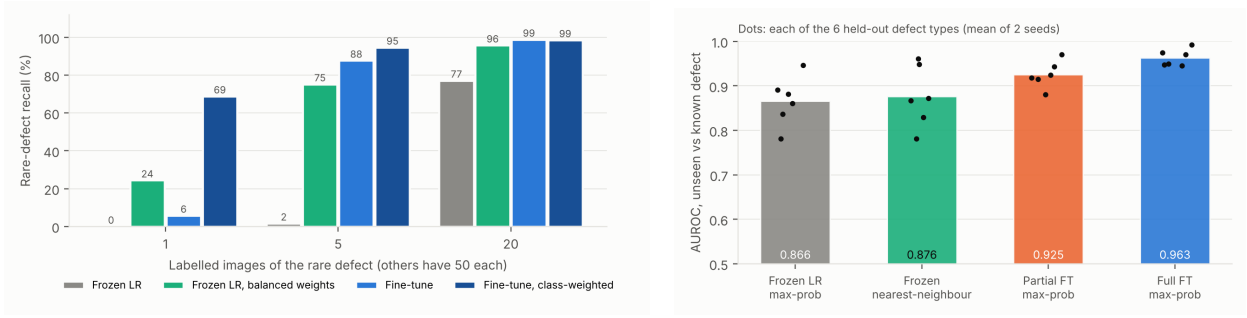
Labels per class k	1	2	5	10	50	240
Frozen CLIP + LR (C=10)	73.6±2.3	82.4±2.6	90.5±2.1	94.4±1.7	98.3±0.9	99.4±0.7
Frozen CLIP + LR (tuned)	74.6±2.3	84.3±1.8	92.7±1.5	96.5±1.5	98.7±0.8	99.4±0.7
Partial fine-tune (CLIP)	82.3±2.9	89.4±2.7	97.3±1.4	98.7±0.8	99.2±0.6	99.9±0.3
Full fine-tune (CLIP)	**87.3±4.7**	**92.4±3.9**	97.1±1.5	**99.4±0.5**	**99.7±0.3**	**100.0±0.0**
Full fine-tune (ImageNet ViT)	79.0±4.0	89.0±1.6	95.6±2.0	98.8±0.4	99.7±0.3	99.9±0.1

$k = 5$ and 1.4 at $k = 20$ (0.2 to 2.5); at $k = 50$ and 240 the intervals include zero. Full fine-tuning reaches 95% accuracy with 3 labels per class and 99% with 10; tuned frozen features need 10 and 240, untuned ($C = 10$) ones 20 and 240. Tuning the frozen classifier lifts its accuracy at $k = 10$ from 94.4% to 96.5%, so an untuned baseline would exaggerate the benefit of fine-tuning. Partial fine-tuning is 5.0 and 3.0 points below full at $k = 1, 2$ (both CIs exclude zero) and indistinguishable from $k = 3$ onward. CLIP beats the ImageNet ViT by 8.4 points at $k = 1$ (5.7 to 11.0), but from $k = 5$ the difference is not reliable. At $k = 1$ the hardest classes for the full model are scratches (recall 0.66), pitting (0.80) and inclusion (0.85); crazing, patches and rolled-in scale exceed 0.95.

Calibration and confident errors. Fine-tuned CLIP is well calibrated (ECE 0.008 at $k = 10$; log-loss 0.025), whereas the frozen classifier at $C = 10$ is under-confident (ECE 0.38; tuning cuts it to 0.07). However, at $k = 1$ the fully fine-tuned model misclassifies 2.8% of test images with probability ≥ 0.9 (22% of its errors), and the ImageNet ViT 8.3% (39% of its errors). The frozen classifier makes no confident errors because it is rarely confident. Average calibration therefore does not guarantee that a high-confidence prediction can be trusted.

Rare defects. With one labelled example of a rare class (others 50), plain fine-tuning finds 6% of its test images and frozen features 0%. Class-weighted fine-tuning finds 69% (52–85% depending on the class) at precision 1.0, and balanced-weight frozen features 24%. With 5 examples plain fine-tuning finds 88% and weighting adds a further 7 points; at 20 examples all fine-tuned variants exceed 98% (Fig. 2a).

Unseen defects. Using maximum softmax probability [13], fine-tuned CLIP separates an unseen defect type from known ones best (AUROC 0.963; false-positive rate at 95% recall 0.12), ahead of partial fine-tuning (0.925), frozen nearest-neighbour distance (0.876) and frozen logistic regression (0.866; Fig. 2b). The spread across defect types and methods is large (0.78–0.99), so a defect-specific alarm threshold would be needed.



(a) Rare-defect recall

(b) Unseen-defect AUROC

Figure 2: (a) Recall for a rare class (mean over 3 classes $\times$ 2 seeds). (b) AUROC, dots are the six held-out defects.

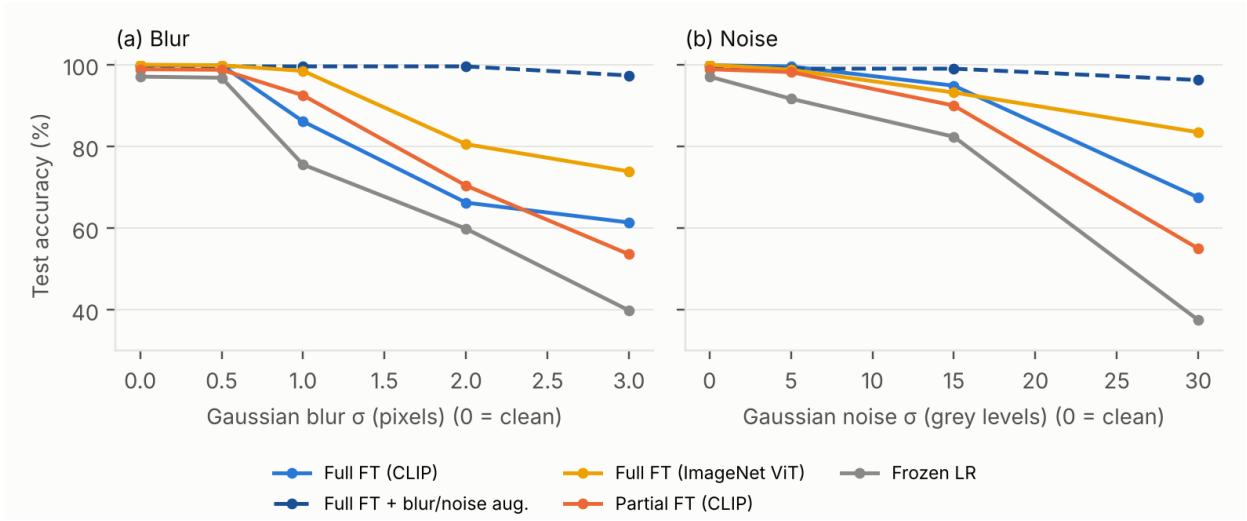


Figure 3: Accuracy under blur and noise ($k = 20$, fold 0, 2 seeds).

Corruptions. Models trained on clean images ($k = 20$) degrade sharply under blur and noise [14] (Fig. 3): at the heaviest blur, accuracy is 61% (CLIP full), 74% (ImageNet ViT) and 40% (frozen); at the heaviest noise, 68%, 84% and 38%. Adding blur and noise augmentation restores 97% and 96%. The augmentation protects only against the distortions it was trained on: under strong brightening ($\times 1.4$) the augmented model is slightly worse (96.4% versus 98.5%; 95.6% versus 99.0% at $k = 240$). These tests use one fold's 360 test images and two seeds.

Sensitivity and zero-shot. Fine-tuned accuracy varies little across three learning rates (96.0–96.9% at $k = 5$; 99.0–99.4% at $k = 20$), while the frozen classifier is very sensitive to regularisation (91.2% to 98.5% across the four values at $k = 20$). Zero-shot CLIP, with prompts chosen on the training pool, scores 19.7% against 16.7% for guessing, with half of all predictions assigned to "patches"; the text encoder has no useful vocabulary for these defects, although the image encoder's features are useful for them.

Qualitative checks. Occlusion maps (Fig. 4) are one image per class and scaled per panel, so they only indicate that models rely on different regions: for the inclusion example the frozen model gives the true class probability 0.20 and the fine-tuned model 0.97. No causal claims are made from them.

Limitations of the evidence. Saturation above 50 labels, a single homogeneous dataset, 360 test images per fold, non-independent folds, and one fold only for the corruption and sensitivity tests limit generality. No convolutional baseline is included, so claims concern ViTs only.

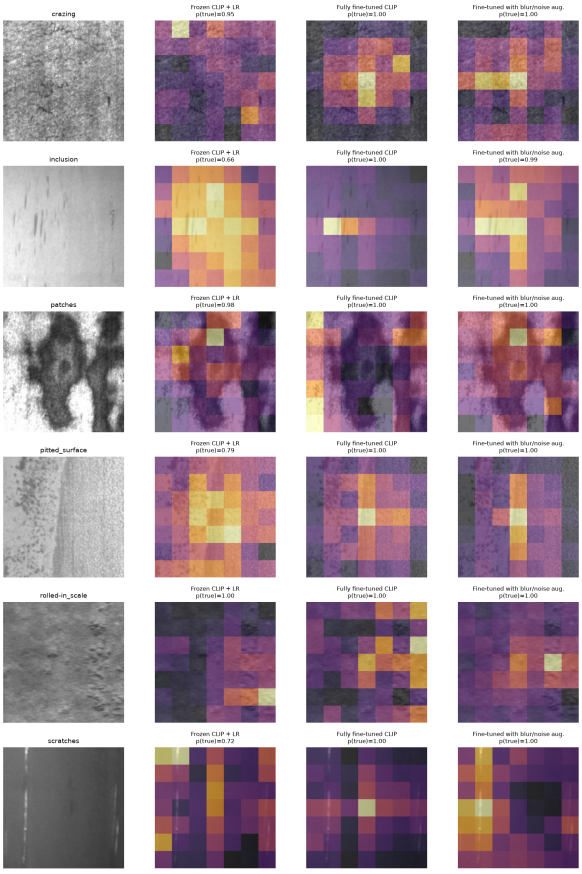


Figure 4: Occlusion sensitivity: brighter areas lower the true-class probability more when hidden.

5. Ethical and Societal Considerations

Safety and accountability. A missed defect can ship a faulty product, and a false alarm wastes steel. The confident-error result means a high softmax score cannot by itself authorise automatic release; human review of low-confidence and novel-looking items, a calibrated threshold per defect type, and logging every decision are recommended. Responsibility for missed defects must stay with the plant's quality function, not the model, in line with the accountability and transparency that Dubai's AI ethics toolkit asks organisations to prioritise [15].

Data, bias and provenance. NEU contains images of steel only; no people appear, so personal-data rules such as the UAE federal data protection law [16] are not engaged. Deploying cameras on a shop floor could change that, and would require lawful processing and data minimisation. CLIP was pretrained on web-scale image–text data whose contents could not be audited [7]; bias there is unlikely to matter for steel textures, but this was not measured. The more relevant bias is dataset bias: one laboratory setup and homogeneous images mean reported accuracy may not transfer to other plants, camera positions or steel grades; models should be revalidated on local data. Documenting the data and models with datasheets and model cards [17, 18] is recommended; the repository includes a short model card.

Workforce. The UAE AI strategy lists manufacturing among sectors with high automation potential and stresses training for new jobs [1]. A model like this one should assist inspectors, who handle low-confidence and novel cases, rather than replace them. Because a few labels give useful accuracy, staff can help label and audit local data, which builds their skills rather than only displacing them.

Environment and transparency. Fine-tuning costs far more compute than extracting frozen features, a footprint that grows with experiment count [19]. The experiment suite was run on a laptop; energy was not

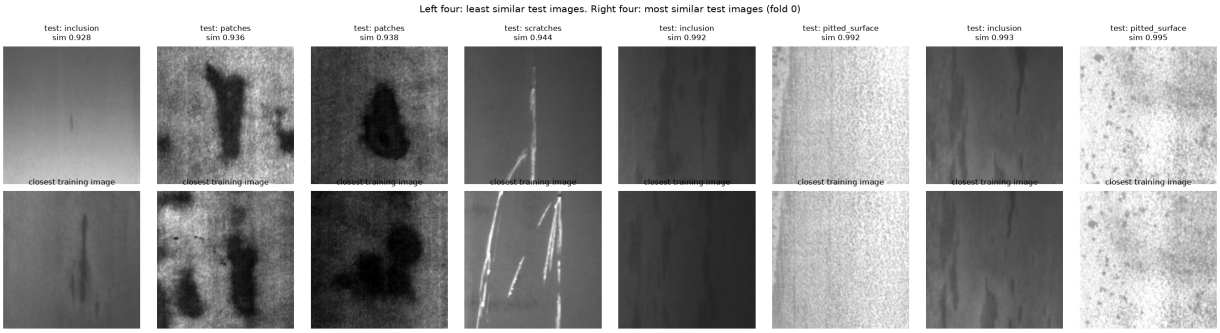


Figure 5: Least similar (left four) and most similar (right four) test images with their closest training image.

measured, and the choice to run many repeated experiments is part of that cost. Frozen features are therefore recommended when labels are plentiful, since they perform within about one point of fine-tuning. Finally, all prompt and parameter variants are reported, including the weak zero-shot result, rather than only favourable settings.

6. Conclusion and Future Work

When labels are very scarce (1 to 5 per class), fully fine-tuning a CLIP ViT is clearly better than a frozen encoder and reaches 95% with 3 labels per class. For rare defects, class-weighted fine-tuning is the best of the methods tested. With 50 or more labels, a tuned frozen encoder is nearly as good at a fraction of the cost. Fine-tuned models are well calibrated on average yet still make confident mistakes, and robustness to blur and noise needs specific training. Limitations include a homogeneous single dataset and few folds for some tests. Future work: a second, more varied dataset (for example GC10-DET or X-SDD); detection rather than classification; parameter-efficient or weight-averaged fine-tuning to protect pretrained features [20]; calibrated or conformal abstention; and testing under a camera or plant shift.

References

[1] UAE Government, “UAE national strategy for artificial intelligence 2031.” https://u.ae/-/media/AI-publications/UAE-National-Strategy-for-Artificial-Intelligence-2031.pdf. Accessed October 2026.

[2] A. Dosovitskiy, L. Beyer, A. Kolesnikov, D. Weissenborn, X. Zhai, T. Unterthiner, M. Dehghani, M. Minderer, G. Heigold, S. Gelly, J. Uszkoreit, and N. Houlsby, “An image is worth 16x16 words: Transformers for image recognition at scale,” in *Proc. International Conference on Learning Representations (ICLR)*, 2021.

[3] K. Song and Y. Yan, “A noise robust method based on completed local binary patterns for hot-rolled steel strip surface defects,” *Applied Surface Science*, vol. 285, pp. 858–864, 2013.

[4] Y. He, K. Song, Q. Meng, and Y. Yan, “An end-to-end steel surface defect detection approach via fusing multiple hierarchical features,” *IEEE Transactions on Instrumentation and Measurement*, vol. 69, no. 4, pp. 1493–1504, 2020.

[5] Kaggle user kaustubhdikshit, “NEU surface defect database (kaggle copy).” https://www.kaggle.com/datasets/kaustubhdikshit/neu-surface-defect-database. Accessed October 2026.

[6] X. Lv, F. Duan, J.-J. Jiang, X. Fu, and L. Gan, “Deep metallic surface defect detection: The new benchmark and detection network,” *Sensors*, vol. 20, no. 6, p. 1562, 2020.

[7] A. Radford, J. W. Kim, C. Hallacy, A. Ramesh, G. Goh, S. Agarwal, G. Sastry, A. Askell, P. Mishkin, J. Clark, G. Krueger, and I. Sutskever, "Learning transferable visual models from natural language supervision," in *Proc. International Conference on Machine Learning (ICML), PMLR 139*, pp. 8748–8763, 2021.

[8] G. Ilharco, M. Wortsman, R. Wightman, C. Gordon, N. Carlini, R. Taori, A. Dave, V. Shankar, H. Namkoong, J. Miller, H. Hajishirzi, A. Farhadi, and L. Schmidt, "OpenCLIP." Software, https://doi.org/10.5281/zenodo.5143773, 2021.

[9] A. Steiner, A. Kolesnikov, X. Zhai, R. Wightman, J. Uszkoreit, and L. Beyer, "How to train your ViT? data, augmentation, and regularization in vision transformers," *Transactions on Machine Learning Research*, 2022.

[10] R. Wightman, "PyTorch image models (timm)." Software, https://doi.org/10.5281/zenodo.4414861, 2019.

[11] H. He and E. A. Garcia, "Learning from imbalanced data," *IEEE Transactions on Knowledge and Data Engineering*, vol. 21, no. 9, pp. 1263–1284, 2009.

[12] C. Guo, G. Pleiss, Y. Sun, and K. Q. Weinberger, "On calibration of modern neural networks," in *Proc. International Conference on Machine Learning (ICML), PMLR 70*, pp. 1321–1330, 2017.

[13] D. Hendrycks and K. Gimpel, "A baseline for detecting misclassified and out-of-distribution examples in neural networks," in *Proc. International Conference on Learning Representations (ICLR)*, 2017.

[14] D. Hendrycks and T. Dietterich, "Benchmarking neural network robustness to common corruptions and perturbations," in *Proc. International Conference on Learning Representations (ICLR)*, 2019.

[15] Gulf News, "Smart dubai launches guidelines on ethical use of artificial intelligence." https://gulfnews.com/technology/smart-dubai-launches-guidelines-on-ethical-use-of-artificial-intelligence-1.61317752, 2019. Accessed October 2026.

[16] Clyde & Co, "UAE issues landmark personal data protection law (federal decree-law no. 45 of 2021)." https://www.clydeco.com/en/insights/2021/11/uae-issues-landmark-personal-data-protection-law, 2021. Accessed October 2026.

[17] T. Gebru, J. Morgenstern, B. Vecchione, J. W. Vaughan, H. Wallach, H. Daumé III, and K. Crawford, "Datasheets for datasets," *Communications of the ACM*, vol. 64, no. 12, pp. 86–92, 2021.

[18] M. Mitchell, S. Wu, A. Zaldivar, P. Barnes, L. Vasserman, B. Hutchinson, E. Spitzer, I. D. Raji, and T. Gebru, "Model cards for model reporting," in *Proc. ACM Conference on Fairness, Accountability, and Transparency (FAT*)*, pp. 220–229, 2019.

[19] E. Strubell, A. Ganesh, and A. McCallum, "Energy and policy considerations for deep learning in NLP," in *Proc. Annual Meeting of the Association for Computational Linguistics (ACL)*, pp. 3645–3650, 2019.

[20] A. Kumar, A. Raghunathan, R. Jones, T. Ma, and P. Liang, "Fine-tuning can distort pretrained features and underperform out-of-distribution," in *Proc. International Conference on Learning Representations (ICLR)*, 2022.